Supporting Research Note

Communicating AI Risk Across Languages

Prepared for an IASEAI'27 talk proposal | 6 October 2026

Purpose

This ongoing project examines how three AI safety-related events are presented in English, Simplified Chinese, and Traditional Chinese video content on Bilibili and YouTube, and how audiences respond beneath those videos. The proposed talk connects this comparative research design with practical questions for AI safety education and communication.

Questions

1. How are societal threats and responses framed in video titles and full narratives?
2. How do audience interactions express perceived threat, attribution, skepticism, and possible responses?
3. To what extent do these supply and reception measures co-vary at the video level?

Corpus and units

The core event corpus contains 4,331 titles across three events. A view-coverage selection contains 304 videos: 35 from Bilibili, 236 English-language YouTube videos, and 33 Chinese-language YouTube videos. Platform counts and language-ecosystem counts are different conventions and must not be treated as interchangeable. Full transcripts are available for the selected videos.

The collected interaction archive contains 6,523 records. Of these, 5,967 link to the selected 304 videos and form the current restricted reception corpus. The remaining 556 interactions belong to videos outside this selection and are excluded from that aligned analysis. Comments and timestamped on-screen messages are distinct formats; neither records a representative sample of viewers.

Planned analysis

Titles and transcripts are coded separately. The framework examines societal threat and efficacy, overall evaluation, proposed responses, and their intended recipients. Reception coding examines perceived threat and attribution or efficacy exits. Contextual close reading supports interpretation and checks whether automated classifications fit the underlying discourse.

Analyses of supply–reception relationships use videos as the alignment unit and account for multiple interactions under the same video. Language, platform, event mix, format, and duration are considered potential sources of confounding. Small cells and unequal coverage limit comparisons.

Status and validation

Data collection and transcription for the selected set are complete. The 5,967 linked interactions have structured recoding records; formal schema and quotation checks have been completed. These checks do not establish semantic accuracy or independent coding reliability. Recoding and descriptive aggregation are now complete for all 304 transcripts (840 windows) and 5,967 linked interactions. Independent reliability assessment, semantic review, and segmentation/off-topic sensitivity remain outstanding. Preliminary descriptive results and executable aggregation scripts are available in this repository.

Statistics from earlier withdrawn outputs are not offered as evidence. The current preliminary results derive from the resolved recoding records; they remain descriptive and unadjusted. Any quantitative results used in the presentation will require traceable item-level records, reproducible aggregation, and explicit reliability and uncertainty reporting. No causal mediation claim is made.

Presentation contribution

The talk will show how a communication audit can distinguish attention-seeking hooks, substantive risk explanations, and situated audience responses. It will offer practical questions about threat, responsibility, and available action. Recommendations will be presented as proposals to test, rather than interventions with established effects.

Limits

The event and view-coverage design is selective. Observed interactions cannot establish population attitudes, geographical residence, or the private beliefs of commenters. Platform and language effects cannot automatically be separated. Statistical associations, if supported, will not establish that a video caused an audience response.

Preliminary descriptive observations

In core transcripts, risk-more coding and institutional recommendations co-occur in 13/21 Simplified Chinese, 10/10 Traditional Chinese and 165/201 English videos. Recommendations principally address industry and political institutions. Comments have technical/neutral discourse as their largest category in all three groups, while commercial-motive skepticism, othering and danger-control demands have different observed proportions. Risk-benefit score and danger-control comment share have an unadjusted Spearman correlation of -0.3193 across 101 comment-bearing videos. This is a descriptive association, not a causal or significance claim. The current manuscript's findings, discussion and conclusion are included as a public excerpt, with source-text case paragraphs omitted.